

Learning algorithms are the seeds, data is the soil, and the learned programs are the grown plants. The machine-learning expert is like a farmer, sowing the seeds, irrigating and fertilizing the soil, and keeping an eye on the health of the crop but otherwise staying out of the way.



Lecture 8

Naïve Bayes & *Review of the Basics*

[Haiping Lu](#) - [MLAI19](#)

Review Preference Poll: 57

Question 1: Multiple Answer

Average

Which part do you want me to review/explain in more detail in Lecture 8?

Correct Answers	Percent Correct
✓ Bayesian regression	54.385%
✓ Pytorch & Deep learning general	45.614%
✓ PCA	56.14%
✓ K-means clustering	42.105%
✓ Autoencoder	47.368%
✓ Convolutional neural network	40.35%
✓ How to run lab 6 and 7 notebooks	19.298%

Week 8 Contents / Objectives

Part A

- Probabilistic Classification
- Naïve Bayes Classifier

Review

- Bayesian Regression
- PCA
- Autoencoder

Week 8 Contents / Objectives

Part A

- **Probabilistic Classification**
- Naïve Bayes Classifier

Review

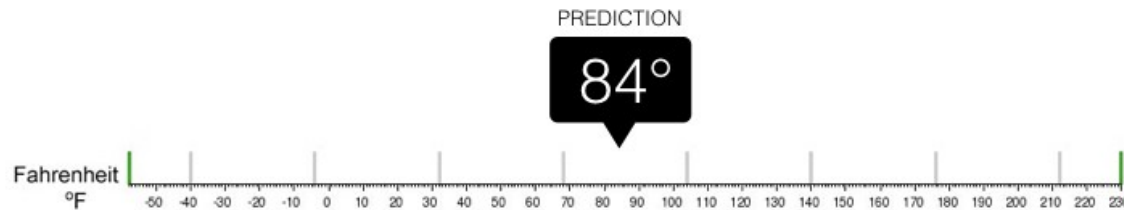
- Bayesian Regression
- PCA
- Autoencoder

Regression vs Classification



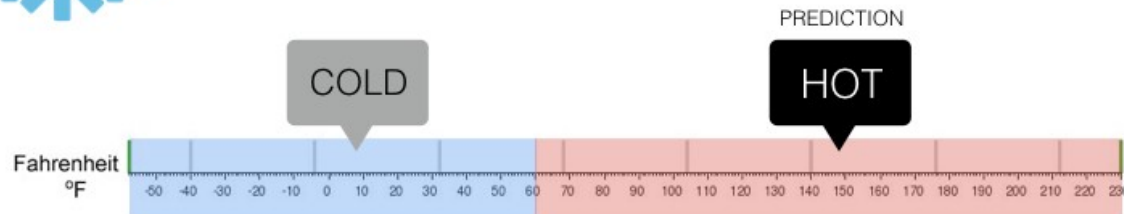
Regression

What is the temperature going to be tomorrow?



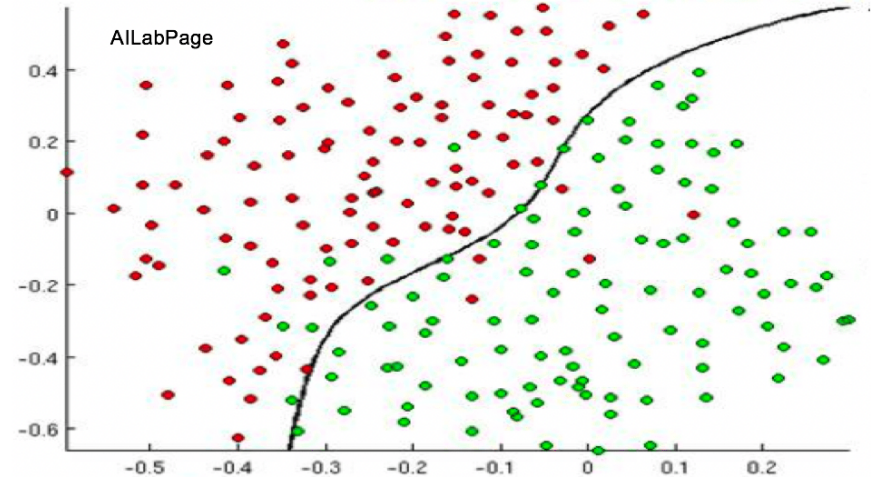

Classification

Will it be Cold or Hot tomorrow?



Source: <https://towardsdatascience.com/regression-or-classification-linear-or-logistic-f093e8757b9c>

Regression vs Classification



Regression

The system attempts to predict a value for an input based on past data.

Example – 1. Temperature for tomorrow



Classification

In classification, predictions are made by classifying them into different categories.

Example – 1. Type of cancer 2. Cancer Y/N

AIlabPage

Probabilistic Classification

- Training classifiers: estimating $f: X \rightarrow Y$, or $P(Y|X)$
- **Discriminative** classifiers
 - Assume some functional form for $P(Y|X)$
 - Estimate parameters of $P(Y|X)$ directly from training data
- **Generative** classifiers
 - Assume some functional form for $P(X|Y)$, $P(X)$
 - Estimate parameters of $P(X|Y)$, $P(X)$ directly from training data
 - Use Bayes rule to calculate $P(Y|X = x_i)$
- **Question:** Is Bayesian regression **discr/gener.?**

Bayes Classifier: MAP

- Training set: p -dimensional data $\mathbf{X} = (x_1, x_2, \dots, x_p)$ and associated class label $y \in \{C_1, C_2, \dots, C_m\}$, i.e., there are m classes. We have n such pairs $(\mathbf{X}_1, y_1), (\mathbf{X}_2, y_2), \dots, (\mathbf{X}_n, y_n)$
- Classification is to derive the maximal $P(C_i|\mathbf{X})$
- Generative classifiers: the Maximum A Posteriori (MAP) derived from Bayes' theorem

$$P(C_i|\mathbf{X}) = \frac{P(\mathbf{X}|C_i)P(C_i)}{P(\mathbf{X})}$$

Question: we can ignore $P(\mathbf{X})$ for classification. Why?

- $P(\mathbf{X})$ is constant for all classes \rightarrow we only need to maximize $P(\mathbf{X}|C_i)P(C_i)$, i.e., compute **1)** prior, and **2)** likelihood.

Week 8 Contents / Objectives

Part A

- Probabilistic Classification
- **Naïve Bayes Classifier**

Review

- Bayesian Regression
- PCA
- Autoencoder

Naïve Bayes Classifier

- Simplified assumption: variables are independent conditioned on the class label

$$P(\mathbf{X}|C_i) = \prod_{k=1}^p P(x_k|C_i) = P(x_1|C_i) \times P(x_2|C_i) \times \cdots \times P(x_p|C_i)$$

- Greatly reduce the #parameters & computational cost
- For categorical variables, we can estimate their probability by *counting its #occurrence divided by the total number of related samples*
- For continuous variables, we can use Gaussian to model them or *convert to categorical via splitting/binning*

Naïve Bayes Example: Data

Split into 3

Class:

C1:buys_computer = 'yes'

C2:buys_computer = 'no'

New data *to be classified*:

X = (age <=30,

Income = medium,

Student = yes

Credit_rating = Fair)

age	income	student	credit_rating	comp
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no

age	income	student	credit_rating	com
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no

Naïve Bayes Example:

- $P(C_i)$: $P(\text{buys_computer} = \text{"yes"}) =$
 $P(\text{buys_computer} = \text{"no"}) =$

- Compute $P(\mathbf{X}|C_i)$ for each class

$$P(\text{age} = \text{"<=30"} \mid \text{buys_computer} = \text{"yes"}) =$$

$$P(\text{income} = \text{"medium"} \mid \text{buys_computer} = \text{"yes"}) =$$

$$P(\text{student} = \text{"yes"} \mid \text{buys_computer} = \text{"yes"}) =$$

$$P(\text{credit_rating} = \text{"fair"} \mid \text{buys_computer} = \text{"yes"}) =$$

Similarly work out those for $\text{buys_computer} = \text{"no"}$

- $\mathbf{X} = (\text{age} \leq 30, \text{income} = \text{medium}, \text{student} = \text{yes}, \text{credit_rating} = \text{fair})$

$$P(\mathbf{X}|C_i) : P(\mathbf{X}|\text{buys_computer} = \text{"yes"}) = 0.222 \times 0.444 \times 0.667 \times 0.667 = 0.044$$

$$P(\mathbf{X}|\text{buys_computer} = \text{"no"}) = 0.6 \times 0.4 \times 0.2 \times 0.4 = 0.019$$

$$P(\mathbf{X}|C_i)P(C_i) : P(\mathbf{X}|\text{buys_computer} = \text{"yes"}) \times P(\text{buys_computer} = \text{"yes"}) = 0.028$$

$$P(\mathbf{X}|\text{buys_computer} = \text{"no"}) \times P(\text{buys_computer} = \text{"no"}) = 0.007$$

Therefore, \mathbf{X} belongs to class ($\text{"buys_computer} = \text{yes"}$)

Another Example for Naïve Bayes

New patient:

P34=M, P61=M, BMI = H

Best guess at cancer field ?

which of these gives the highest value?

P34 level	P61 level	BMI	Prostate cancer
High	Low	Medium	Y
Medium	Low	Medium	Y
Low	Low	High	Y
Low	High	Low	N
Low	Low	Low	N
Medium	Medium	Low	N
High	Low	Medium	Y
High	Medium	Low	N
Low	Low	High	N
Medium	High	High	Y

$$P(p34=M | Y) \times P(p61=M | Y) \times P(BMI=H | Y) \times P(\text{cancer} = Y)$$

$$P(p34=M | N) \times P(p61=M | N) \times P(BMI=H | N) \times P(\text{cancer} = N)$$

Another Example for Naïve Bayes

New patient:

P34=M, P61=M, BMI = H

Best guess at cancer field ?

which of these gives the highest value?

P34 level	P61 level	BMI	Prostate cancer
High	Low	Medium	Y
Medium	Low	Medium	Y
Low	Low	High	Y
Low	High	Low	N
Low	Low	Low	N
Medium	Medium	Low	N
High	Low	Medium	Y
High	Medium	Low	N
Low	Low	High	N
Medium	High	High	Y

$$P(\mathbf{p34=M} \mid \mathbf{Y}) \times P(\mathbf{p61=M} \mid \mathbf{Y}) \times P(\mathbf{BMI=H} \mid \mathbf{Y}) \times P(\mathbf{cancer = Y})$$

$$P(\mathbf{p34=M} \mid \mathbf{N}) \times P(\mathbf{p61=M} \mid \mathbf{N}) \times P(\mathbf{BMI=H} \mid \mathbf{N}) \times P(\mathbf{cancer = N})$$

Another Example for Naïve Bayes

New patient:

P34=M, P61=M, BMI = H

Best guess at cancer field ?

which of these gives the highest value?

P34 level	P61 level	BMI	Prostate cancer
High	Low	Medium	Y
Medium	Low	Medium	Y
Low	Low	High	Y
Low	High	Low	N
Low	Low	Low	N
Medium	Medium	Low	N
High	Low	Medium	Y
High	Medium	Low	N
Low	Low	High	N
Medium	High	High	Y

$$P(p34=M | Y) \times \textcolor{red}{P(p61=M | Y)} \times P(BMI=H | Y) \times P(\text{cancer} = Y)$$

$$P(p34=M | N) \times P(p61=M | N) \times P(BMI=H | N) \times P(\text{cancer} = N)$$

Another Example for Naïve Bayes

New patient:

P34=M, P61=M, BMI = H

Best guess at cancer field ?

which of these gives the highest value?

P34 level	P61 level	BMI	Prostate cancer
High	Low	Medium	Y
Medium	Low	Medium	Y
Low	Low	High	Y
Low	High	Low	N
Low	Low	Low	N
Medium	Medium	Low	N
High	Low	Medium	Y
High	Medium	Low	N
Low	Low	High	N
Medium	High	High	Y

$$P(p34=M | Y) \times P(p61=M | Y) \times \mathbf{P(BMI=H | Y)} \times P(\text{cancer} = Y)$$

$$P(p34=M | N) \times P(p61=M | N) \times P(BMI=H | N) \times P(\text{cancer} = N)$$

Another Example for Naïve Bayes

New patient:

P34=M, P61=M, BMI = H

Best guess at cancer field ?

which of these gives the highest value?

P34 level	P61 level	BMI	Prostate cancer
High	Low	Medium	Y
Medium	Low	Medium	Y
Low	Low	High	Y
Low	High	Low	N
Low	Low	Low	N
Medium	Medium	Low	N
High	Low	Medium	Y
High	Medium	Low	N
Low	Low	High	N
Medium	High	High	Y

$$P(p34=M | Y) \times P(p61=M | Y) \times P(BMI=H | Y) \times \mathbf{P(cancer = Y)}$$

$$P(p34=M | N) \times P(p61=M | N) \times P(BMI=H | N) \times P(cancer = N)$$

Another Example for Naïve Bayes

New patient:

P34=M, P61=M, BMI = H

Best guess at cancer field ?

which of these gives the highest value?

P34 level	P61 level	BMI	Prostate cancer
High	Low	Medium	Y
Medium	Low	Medium	Y
Low	Low	High	Y
Low	High	Low	N
Low	Low	Low	N
Medium	Medium	Low	N
High	Low	Medium	Y
High	Medium	Low	N
Low	Low	High	N
Medium	High	High	Y

$$\begin{array}{llll}
 0.4 & \times 0 & \times 0.4 & \times 0.5 = 0 \\
 0.2 & \times 0.4 & \times 0.2 & \times 0.5 = 0.008
 \end{array}$$

In practice, we finesse the zeroes and use logs:
 (note: $\log(A \times B \times C \times D \times \dots) = \log(A) + \log(B) + \dots$)

New patient:

P34=M, P61=M, BMI = H

Best guess at cancer field ?

which of these gives the
highest value?

P34 level	P61 level	BMI	Prostate cancer
High	Low	Medium	Y
Medium	Low	Medium	Y
Low	Low	High	Y
Low	High	Low	N
Low	Low	Low	N
Medium	Medium	Low	N
High	Low	Medium	Y
High	Medium	Low	N
Low	Low	High	N
Medium	High	High	Y

$\log(0.4)$	$+ \log(0.001)$	$+ \log(0.4)$	$+ \log(0.5) = -4.09$
$\log(0.2)$	$+ \log(0.4)$	$+ \log(0.2)$	$+ \log(0.5) = -2.09$

Numerical Stability

$$\log(a \times b \times c \times \dots) = \log(a) + \log(b) + \log(c) + \dots$$

This helps us to avoid/reduce the **underflow** errors, which we would otherwise get with when multiplying many probabilities, e.g.

$$0.003 \times 0.000296 \times 0.001 \times \dots[\mathbf{100} \text{ fields}] \times 0.042 \dots$$

Avoiding the Zero-Probability

- One conditional prob. = zero \rightarrow predicted prob. = zero
 - Not desired
- Use **Laplacian smoothing**
 - E.g., a dataset with 1000 samples, income=low (0), income=medium (990), and income = high (10)
 - *Adding 1 to each case*
 - $\text{Prob}(\text{income} = \text{low}) = 1/1003$
 - $\text{Prob}(\text{income} = \text{medium}) = 991/1003$
 - $\text{Prob}(\text{income} = \text{high}) = 11/1003$
 - The “smoothed” prob. estimates are close to their “unsmoothed” counterparts

Summary on Naïve Bayes

- Assumption: features independent conditioned on class label
- Advantages
 - Easy to implement and fast to compute
 - Good results obtained in many high-dim cases
- Disadvantages
 - Often dependencies exist among variables, e.g., pixels in image, ...
 - **Question:** *what method* can help reduce the dependency?
- Verdict: A good baseline to start with

Week 8 Contents / Objectives

Part A

- Probabilistic Classification
- Naïve Bayes Classifier

Review

- Bayesian Regression
- PCA
- Autoencoder

Factors for deciding review topics

- Poll
- Mock quiz 2 results
- Discussion board Q&A
- Q&A in Lab

Question 1: Multiple Answer

Average

Which part do you want me to review/explain in more detail in Lecture 8?

Correct Answers	Percent Correct
✓ Bayesian regression	54.385%
✓ Pytorch & Deep learning general	45.614%
✓ PCA	56.14%
✓ K-means clustering	42.105%
✓ Autoencoder	47.368%
✓ Convolutional neural network	40.35%
✓ How to run lab 6 and 7 notebooks	19.298%

Machine Learning Ingredients

- **Data:** +pre-processing (& visualisation), e.g., $\mathcal{N}(0,1)$
- **Model**
 - Structure ~ Architecture \leftarrow expert knowledge
 - Must **specify** before ML, can optimise via cross validation (CV)
 - **Hyper-parameter**, e.g., prior, #degree, layer \leftarrow knowledge
 - Must **specify** (choices) and can optimise via CV (*tuning*)
 - Parameters (theta)
 - Compute/learn parameter, e.g., **weights**, bias \leftarrow optimisation alg.
- Evaluation metric (what's best): loss/error function
- Optimisation: (how to find the best) learnable parameters

Week 8 Contents / Objectives

Part A

- Probabilistic Classification
- Naïve Bayes Classifier

Review

- **Bayesian Regression**
- PCA
- Autoencoder

Bayesian Regression vs Non-Bayes

$$p(\theta | y) \propto p(y | \theta) p(\theta)$$

posterior \propto likelihood \cdot prior

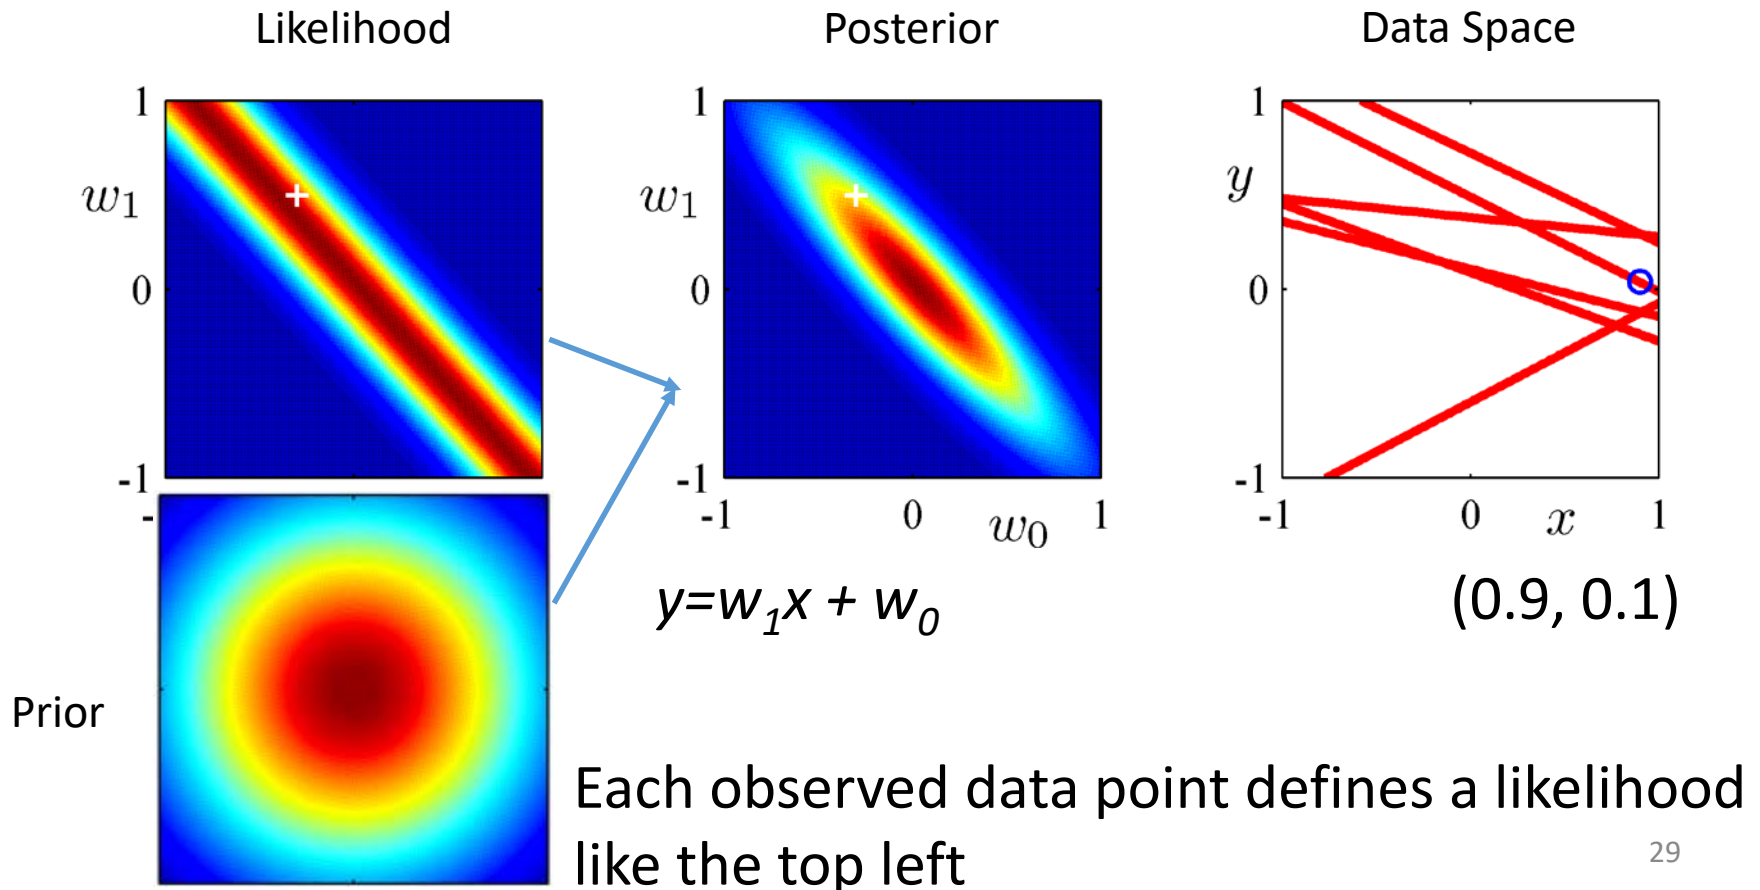
- Bayesian regression
 - Model structure: assuming given in this module
 - Place/specify prior on model parameters (weights)
 - Parameter for prior \rightarrow Hyper-parameters (given or CV)
 - Compute likelihood after observing data
 - Update posterior (\rightarrow new prior) and iterate
 - Metric: MSE, Max A Posterior (MAP); optim: closed/SGD
- Non-Bayesian
 - Model structure: assuming given
 - Compute likelihood after observing data
 - Metric: MSE, Max Likelihood estimation (MLE); optim: ...

Bayesian Regression Ingredients

- **Data:** +pre-processing, e.g., $\mathcal{N}(0,1)$
- **Model**
 - Structure/Architecture: basis function chosen, e.g., poly
 - **Hyper-parameter:** basis function (e.g., degree) & prior hyper
 - Parameters (theta): weights and bias
- Evaluation metric (what's best): MSE
- Optimisation: (how to find the best): closed form for Gaussian distributions, SGD etc. otherwise

Bayesian Regression (S14 of L6)


1 data point observed \rightarrow soft constraint. This
posterior \rightarrow prior for the next data point observed)




Bayesian Regression (S15 of L6)

Another data point observed


Likelihood



Posterior



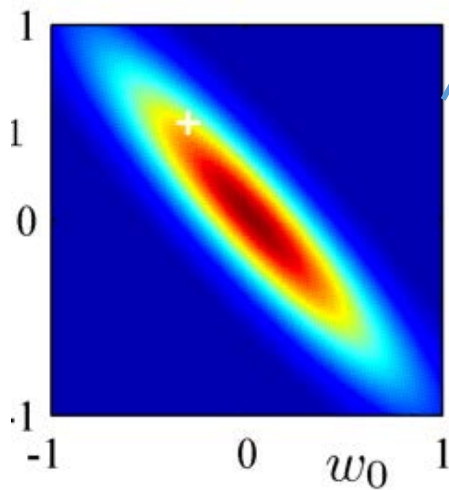
Data Space



$$y = w_1 x + w_0$$

$(-0.7, -0.8)$

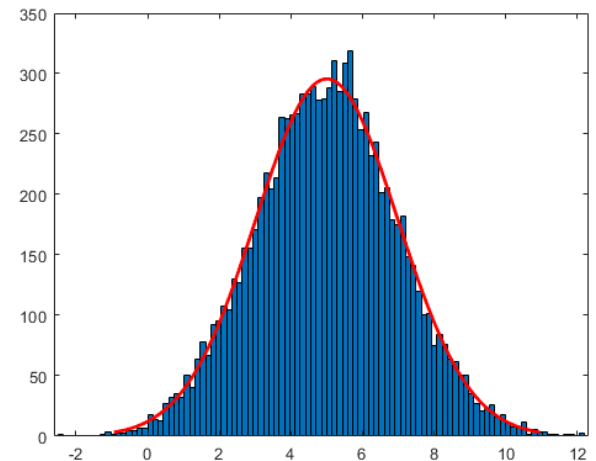


Current
Prior =
Previous
posterior



Question:
What will be the MLE
solution at this point?

Bayesian Regression: Computation

- Gaussian/Normal distribution
 - Knowing the **mean and (co)variance** (std) is sufficient to specify the distribution (*sufficient statistics*)
 - Closed form solution often feasible
 - Solution
 - Density estimation: estimating mean and (co)variance
 - Optimisation: take the mode (max) \rightarrow mean



Lab 6 – Main Trick

$$y = mx + c + \epsilon$$

$$c \sim \mathcal{N}(0, \alpha_1)$$

$$\epsilon \sim \mathcal{N}(0, \sigma^2)$$

$$p(c) = \frac{1}{\sqrt{2\pi\alpha_1}} \exp\left(-\frac{1}{2\alpha_1}c^2\right)$$

$$p(\mathbf{y}|\mathbf{x}, c, m, \sigma^2) = \frac{1}{(2\pi\sigma^2)^{\frac{n}{2}}} \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - mx_i - c)^2\right)$$

$$p(c|\mathbf{y}, \mathbf{x}, m, \sigma^2) = \frac{p(\mathbf{y}|\mathbf{x}, c, m, \sigma^2)p(c)}{p(\mathbf{y}|\mathbf{x}, m, \sigma^2)} = \frac{p(\mathbf{y}|\mathbf{x}, c, m, \sigma^2)p(c)}{\int p(\mathbf{y}|\mathbf{x}, c, m, \sigma^2)p(c)dc}$$

$$= \frac{1}{\sqrt{(2\pi\tau^2)}} \exp\left(-\frac{1}{2\tau^2}(c - \mu)^2\right) \quad p(c|\mathbf{y}, \mathbf{x}, m, \sigma^2) \propto p(\mathbf{y}|\mathbf{x}, c, m, \sigma^2)p(c)$$

$$\log p(c|\mathbf{y}, \mathbf{x}, m, \sigma^2) = -\frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - c - mx_i)^2 - \frac{1}{2\alpha_1}c^2 + \text{const}$$

$$= -\frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - mx_i)^2 - \left(\frac{n}{2\sigma^2} + \frac{1}{2\alpha_1}\right)c^2 + c \frac{\sum_{i=1}^n (y_i - mx_i)}{\sigma^2},$$

$$\log p(c|\mathbf{y}, \mathbf{x}, m, \sigma^2) = -\frac{1}{2\tau^2}(c - \mu)^2 + \text{const},$$

$$\tau^2 = (n\sigma^{-2} + \alpha_1^{-1})^{-1} \quad \mu = \frac{\tau^2}{\sigma^2} \sum_{i=1}^n (y_i - mx_i)$$

Bayesian Regression: Benefits

- Key benefits
 - Enable taking prior knowledge into account
 - Provide uncertainty estimation, predicting an output distribution with mean and **variance**
- Limitation
 - If prior is wrong, ...
 - Complexity
- More depth (derivations for multivariate, etc.)?
<https://www.youtube.com/watch?v=dtkGq9tdYcI>

ML 10.1-10.7 from

<https://www.youtube.com/watch?v=yDLKJtOVx5c&list=PLD0F06AA0D2E8FFBA>

Week 8 Contents / Objectives

Part A

- Probabilistic Classification
- Naïve Bayes Classifier

Review

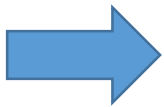
- Bayesian Regression
- **PCA**
- Autoencoder

PCA Ingredients

- **Data:** +pre-processing, e.g., $\mathcal{N}(0, I)$
- **Model**
 - Structure/Architecture: linear projection $\mathbf{y} = \mathbf{U}^T \mathbf{x}$
 - **Hyper-parameter:** lower dimension k
 - Parameters (theta): \mathbf{U}
- Evaluation metric (what's best): max variance
- Optimisation: (how to find the best) eigen-decomposition

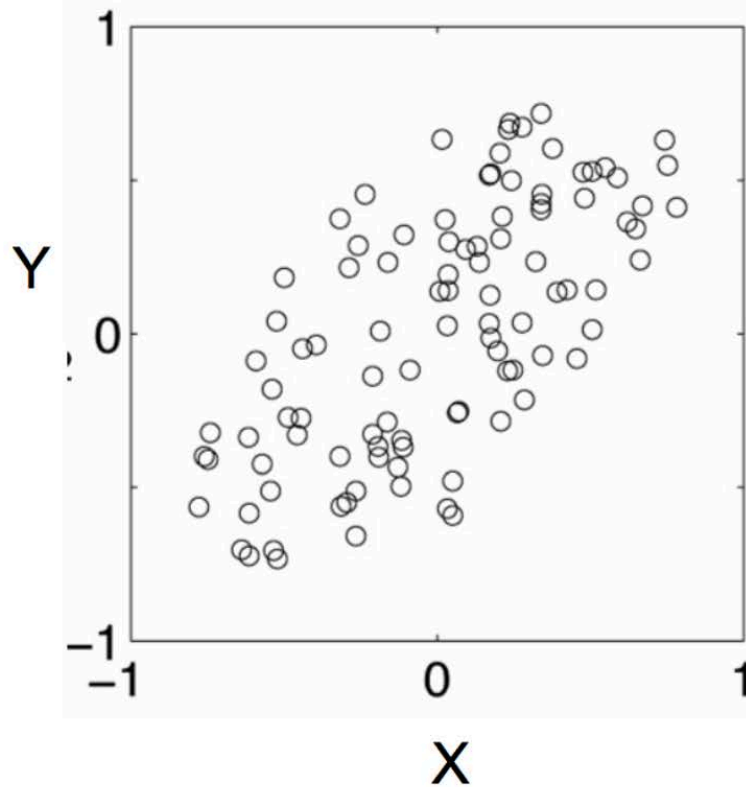
Covariance

- Variance and Covariance:
 - Measure of the “spread” of a set of points around their *center of mass* (mean) (→ similar measurement as *k*-means)
- Variance (**scalar**):
 - Measure of the deviation from the mean for points in **one dimension**
- Covariance (**matrix**):
 - Measure of how much each of the dimensions vary from the mean with **respect to each other (~ correlations)**



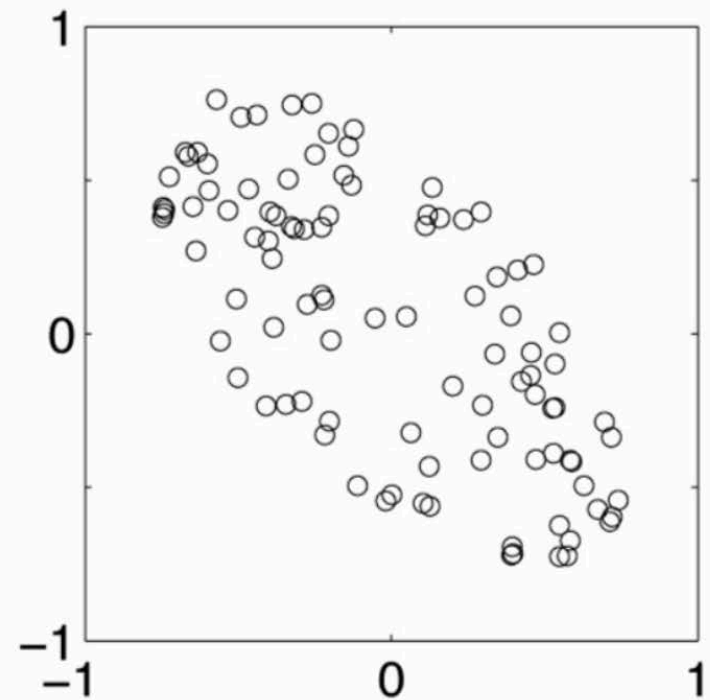
- Covariance is measured between two dimensions
- Covariance sees if there is a relation between the two dimensions
- Covariance between one dimension is the variance

positive covariance



Positive: Both dimensions increase or decrease together

negative covariance



Negative: While one increase the other decrease

Covariance

- Used to find relationships between dimensions in high dimensional data sets
- Scatter matrix: sample-based estimation of covariance matrix

$$q_{jk} = \frac{1}{N} \sum_{i=1}^N (X_{ij} - E(X_j)) (X_{ik} - E(X_k))$$



The Sample mean

- Uncorrelated variables \rightarrow covariance = 0
- Diagonal cov mat \rightarrow all variables are uncorrelated

Form correlated from original by rotating the data space using [rotation matrix](#) \mathbf{R} .

$$p(\mathbf{y}) = \frac{1}{|2\pi\mathbf{D}|^{\frac{1}{2}}} \exp\left(-\frac{1}{2}(\mathbf{y} - \boldsymbol{\mu})^\top \mathbf{D}^{-1}(\mathbf{y} - \boldsymbol{\mu})\right)$$

Multiply the data by a rotation matrix \mathbf{R}

$$p(\mathbf{y}) = \frac{1}{|2\pi\mathbf{D}|^{\frac{1}{2}}} \exp\left(-\frac{1}{2}(\mathbf{R}^\top \mathbf{y} - \mathbf{R}^\top \boldsymbol{\mu})^\top \mathbf{D}^{-1}(\mathbf{R}^\top \mathbf{y} - \mathbf{R}^\top \boldsymbol{\mu})\right)$$

Collect \mathbf{R} to the left and right of \mathbf{D}

$$p(\mathbf{y}) = \frac{1}{|2\pi\mathbf{D}|^{\frac{1}{2}}} \exp\left(-\frac{1}{2}(\mathbf{y} - \boldsymbol{\mu})^\top \mathbf{R} \mathbf{D}^{-1} \mathbf{R}^\top (\mathbf{y} - \boldsymbol{\mu})\right)$$

Let

$$\mathbf{C}^{-1} = \mathbf{R} \mathbf{D}^{-1} \mathbf{R}^\top$$

Rewritten in typical Guassian form

$$p(\mathbf{y}) = \frac{1}{|2\pi\mathbf{C}|^{\frac{1}{2}}} \exp\left(-\frac{1}{2}(\mathbf{y} - \boldsymbol{\mu})^\top \mathbf{C}^{-1}(\mathbf{y} - \boldsymbol{\mu})\right)$$

This gives a **covariance matrix** for correlated variables (the general case):

$$\mathbf{C} = \mathbf{R} \mathbf{D} \mathbf{R}^\top$$

Note $|\mathbf{C}| = |\mathbf{D}|$, see [Determinant of Matrix Product](#)

PCA

Input: $\mathbf{x} \in \mathbb{R}^D: \mathcal{D} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$

Set of basis vectors: $\mathbf{u}_1, \dots, \mathbf{u}_K$

Summarize a D dimensional vector \mathbf{x} with K dimensional feature vector $h(\mathbf{x})$

$$h(\mathbf{x}) = \begin{bmatrix} \mathbf{u}_1 \cdot \mathbf{x} \\ \mathbf{u}_2 \cdot \mathbf{x} \\ \dots \\ \mathbf{u}_K \cdot \mathbf{x} \end{bmatrix}$$

$$\mathbf{U} = [\mathbf{u}_1, \dots, \mathbf{u}_K]$$

Basis vectors are orthonormal

$$\mathbf{u}_i^T \mathbf{u}_j = 0$$

$$||\mathbf{u}_j|| = 1$$

Lagrangian for PCA Solution

- Scatter mat for I/P $\mathbf{S} = \frac{1}{n} \sum_{i=1}^n \left(\mathbf{x}^{(i)} - \boldsymbol{\mu} \right) \left(\mathbf{x}^{(i)} - \boldsymbol{\mu} \right)^{\top}$
- **Question:** what is the scatter mat for the projections?
- Find the first direction via (unit-norm) constrained optimisation, using [Lagrange multipliers](#):

$$L(\mathbf{u}_1, \lambda_1) = \mathbf{u}_1^{\top} \mathbf{S} \mathbf{u}_1 + \lambda_1 (1 - \mathbf{u}_1^{\top} \mathbf{u}_1)$$

- Gradient w.r.t. \mathbf{u}_1 : $\frac{dL(\mathbf{u}_1, \lambda_1)}{d\mathbf{u}_1} = 2\mathbf{S}\mathbf{u}_1 - 2\lambda_1\mathbf{u}_1$
- Set to 0 and rearrange: $\mathbf{S}\mathbf{u}_1 = \lambda_1\mathbf{u}_1$
 - [Eigenvalue problem](#), physical meanings of eigenvalues

Week 8 Contents / Objectives

Part A

- Probabilistic Classification
- Naïve Bayes Classifier

Review

- Bayesian Regression
- PCA
- **Autoencoder**

Convolutional Autoencoder

```
class Autoencoder(nn.Module):
    def __init__(self):
        super(Autoencoder, self).__init__()
        self.encoder = nn.Sequential(
            # 1 input image channel, 16 output channel, 3x3 square convolution
            nn.Conv2d(1, 16, 3, stride=2, padding=1),
            nn.ReLU(),
            nn.Conv2d(16, 32, 3, stride=2, padding=1),
            nn.ReLU(),
            nn.Conv2d(32, 64, 7)
        )
        self.decoder = nn.Sequential(
            nn.ConvTranspose2d(64, 32, 7),
            nn.ReLU(),
            nn.ConvTranspose2d(32, 16, 3, stride=2, padding=1, output_padding=1),
            nn.ReLU(),
            nn.ConvTranspose2d(16, 1, 3, stride=2, padding=1, output_padding=1),
            nn.Sigmoid() #to range [0, 1]
        )

    def forward(self, x):
        x = self.encoder(x)
        x = self.decoder(x)
        return x
```

Autoencoder Ingredients

- **Data:** +pre-processing, e.g., $\mathcal{N}(0,1)$
- **Model**
 - Structure/Architecture: layers defined in nn.module
 - **Hyper-parameter:** layer specs, e.g., #channels, kernel size
 - Parameters (theta): layer weights and biases
- Evaluation metric (what's best): MSE or other
- Optimisation: (how to find the best) backprop

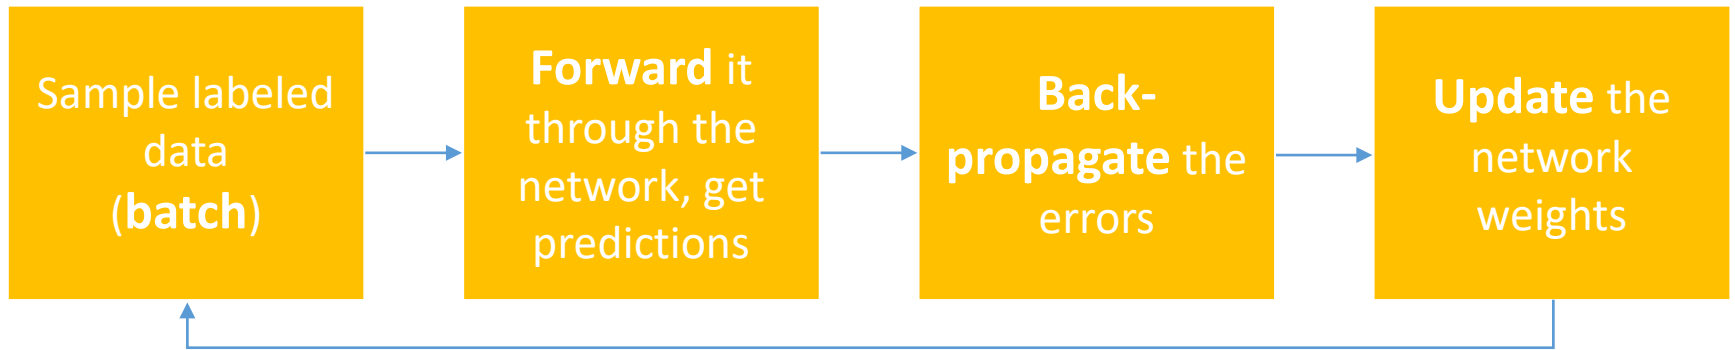
Autoencoder Training

```
#Hyperparameters for training
batch_size=64
learning_rate=1e-3
max_epochs = 20

#Set the random seed for reproducibility
torch.manual_seed(509)
#Choose mean square error loss
criterion = nn.MSELoss()
#Choose the Adam optimiser
optimizer = torch.optim.Adam(myAE.parameters(), lr=learning_rate, weight_decay=1e-5)
#Specify how the data will be loaded in batches (with random shuffling)
train_loader = torch.utils.data.DataLoader(mnist_data, batch_size=batch_size, shuffle=True)
#Storage
outputs = []

#Start training
for epoch in range(max_epochs):
    for data in train_loader:
        img, label = data
        optimizer.zero_grad()
        recon = myAE(img)
        loss = criterion(recon, img)
        loss.backward()
        optimizer.step()
    if (epoch % 3) == 0:
        print('Epoch:{}, Loss:{:.4f}'.format(epoch+1, float(loss)))
    outputs.append((epoch, img, recon),)
```

Training



Data → Model → Metric → Optimisation

Machine Learning Ingredients

- **Data:** +pre-processing (& visualisation), e.g., $\mathcal{N}(0,1)$
- **Model**
 - Structure ~ Architecture \leftarrow expert knowledge
 - Must **specify** before ML, can optimise via cross validation (CV)
 - **Hyper-parameter**, e.g., prior, #degree, layer \leftarrow knowledge
 - Must **specify** (choices) and can optimise via CV (*tuning*)
 - Parameters (theta)
 - Compute/learn parameter, e.g., **weights**, bias \leftarrow optimisation alg.
- Evaluation metric (what's best): loss/error function
- Optimisation: (how to find the best) learnable parameters

Acknowledgement

- Part A used materials from: Jonathan Huang, David Corne, Tom Mitchell, Jiawei Han, Micheline Kamber, and Jian Pei
- Review used additional materials (addition to Lecture 6/7 materials) from: Fereshteh Sadeghi